

DEVOPS: SOFTWARE ARCHITECTURE LAB

PRACTICAL FILE

Bachelor of technology

Information Technology

April, 2022



Submitted by

Parvesh Bhatt

University Roll No. 1905375

Class Roll no. 1921078

Submitted to

Prof. Harjot Kaur Gill

GURU NANAK DEV ENGINEERING COLLEGE
LUDHIANA-141006, INDIA

Contents

1	PRACTICAL 1	1
1.1	Checking for git	1
1.2	Install git on Windows	1
1.3	Install Git on Linux	4
2	PRACTICAL 2	5
3	PRACTICAL 3	6
4	PRACTICAL 4	8
4.1	CREATE BRANCH	8
4.2	MERGE BRANCH	8
4.3	DELETE BRANCH	8
5	PRACTICAL 5	10
6	PRACTICAL 6	13
7	PRACTICAL 7	16
8	PRACTICAL 8	18
9	PRACTICAL 9	21
10	PRACTICAL 10	23
11	PRACTICAL 11	25

1 PRACTICAL 1

INSTALLATION OF GIT

1.1 Checking for git

To see if you already have Git installed, open up your terminal application

- If you're on a Mac, look for a command prompt application called "Terminal".
- If you're on a Windows machine, open the windows command prompt or "Git Bash". [Figure 4]

Once you've opened your terminal application, type `git version`. The output will either tell you which version of Git is installed, or it will alert you that git is an unknown command. If it's an unknown command, read further and find out how to install Git.

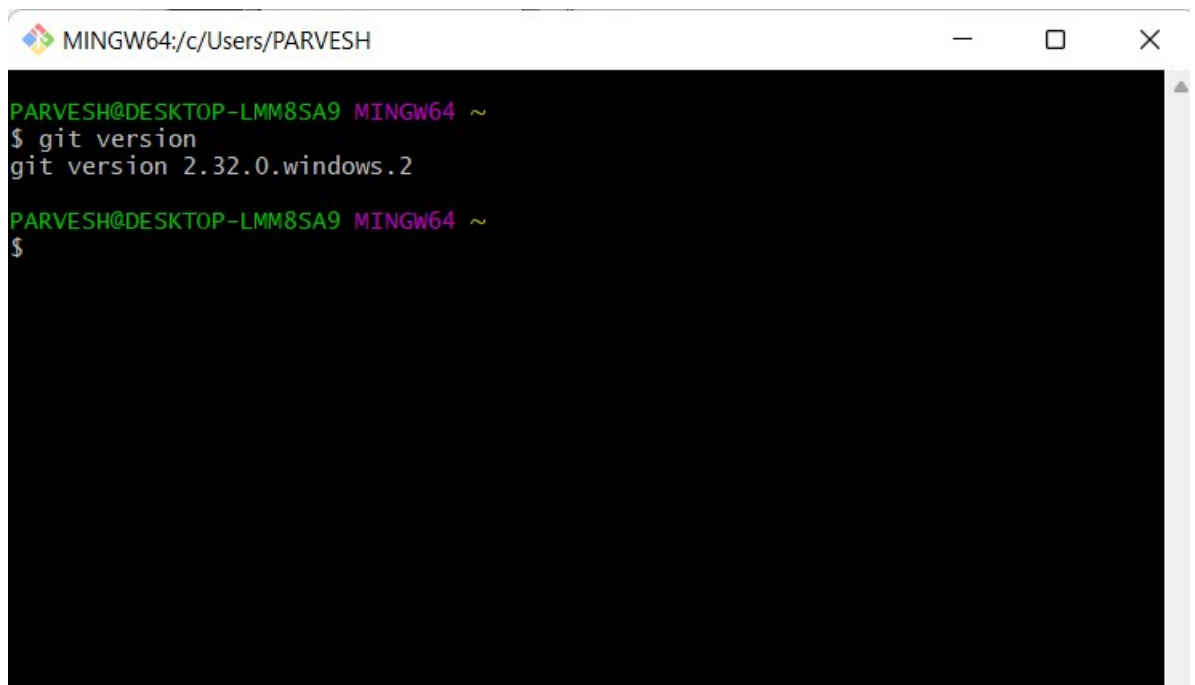
A screenshot of a Windows terminal window titled "MINGW64:/c/Users/PARVESH". The terminal shows a command prompt with the text "PARVESH@DESKTOP-LMM8SA9 MINGW64 ~". The user enters the command "\$ git version", and the terminal outputs "git version 2.32.0.windows.2". The prompt then returns to "\$".

Figure 1: Checking git version

1.2 Install git on Windows

1. Browse to the official Git website: <https://git-scm.com/downloads>

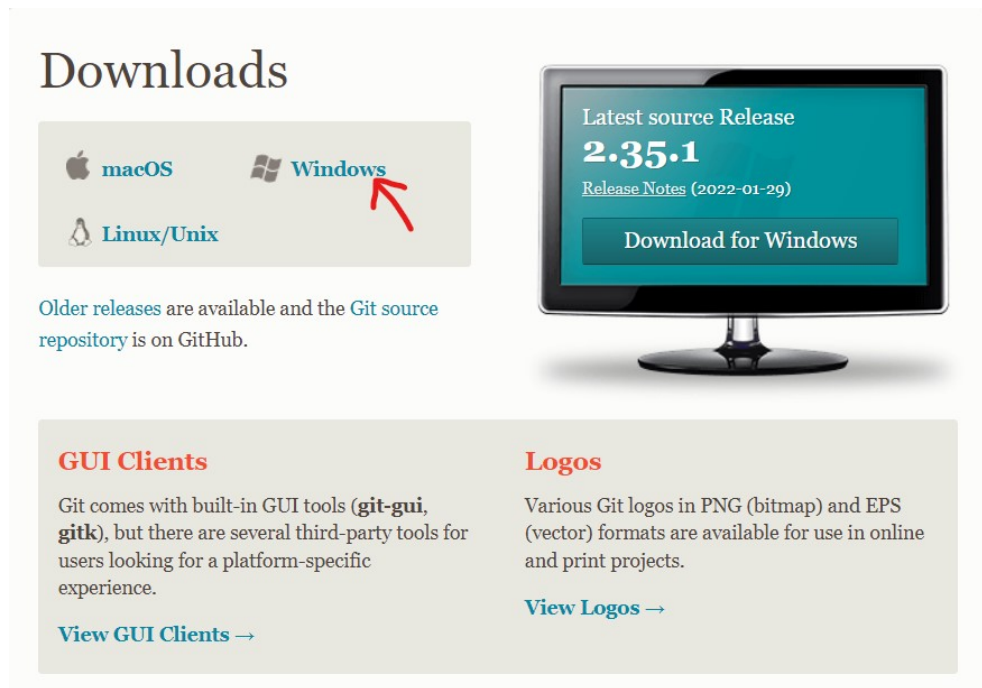


Figure 2: Install git site

2. Click the download link for Windows and allow the download to complete.
3. Browse to the download location (or use the download shortcut in your browser). Double-click the file to extract and launch the installer.
4. Allow the app to make changes to your device by clicking Yes on the User Account Control dialog that opens.

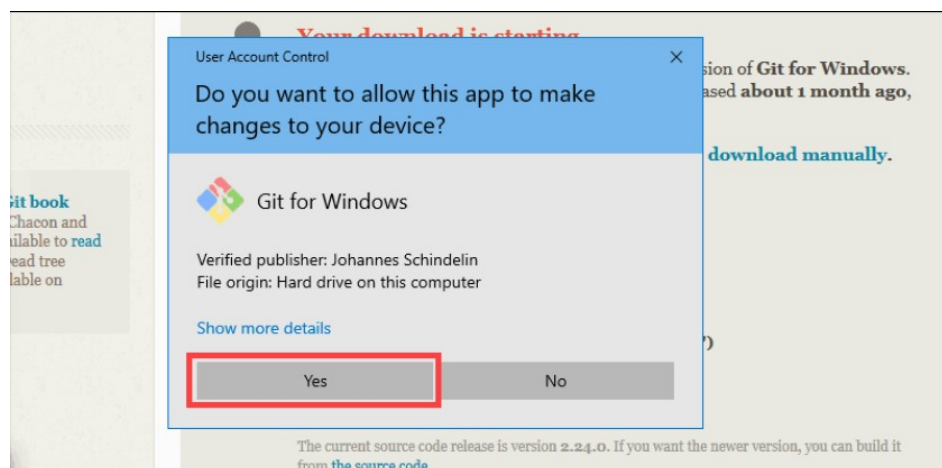


Figure 3: Install git site

5. Review the GNU General Public License, and when you're ready to install, click Next.
6. The installer will ask you for an installation location. Leave the default, unless you have reason to change it, and click Next.

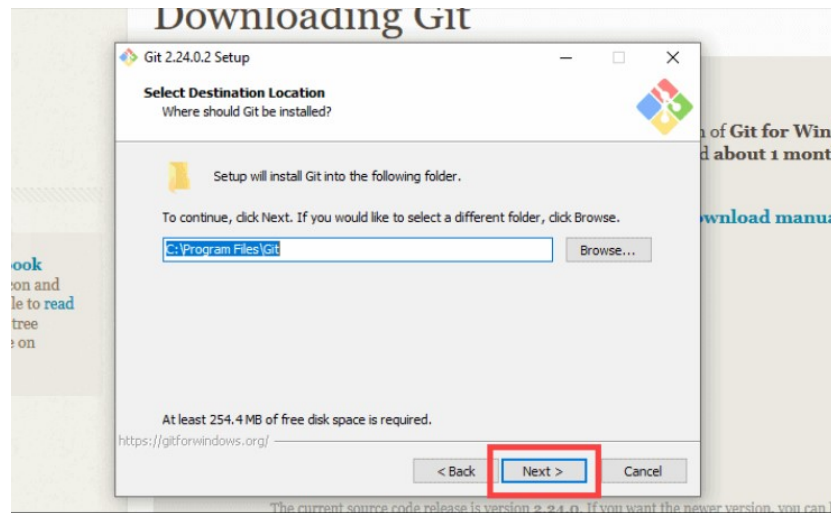


Figure 4: Select install location

7. A component selection screen will appear. Leave the defaults unless you have a specific need to change them and click Next.
8. The installer will offer to create a start menu folder. Simply click Next.
9. Select a text editor you'd like to use with Git. Use the drop-down menu to select Notepad++ (or whichever text editor you prefer) and click Next.
10. The next step allows you to choose a different name for your initial branch. The default is 'master.' Unless you're working in a team that requires a different name, leave the default option and click Next.
11. This installation step allows you to change the PATH environment. The PATH is the default set of directories included when you run a command from the command line. Leave this on the middle (recommended) selection and click Next.

Server Certificates, Line Endings and Terminal Emulators

12. The installer now asks which SSH client you want Git to use. Git already comes with its own SSH client, so if you don't need a specific one, leave the default option and click Next.
13. The next option relates to server certificates. Most users should use the default. If you're working in an Active Directory environment, you may need to switch to Windows Store certificates. Click Next.
14. The next selection converts line endings. It is recommended that you leave the default selection. This relates to the way data is formatted and changing this option may cause problems. Click Next.
15. Choose the terminal emulator you want to use. The default MinTTY is recommended, for its features. Click Next.
16. The installer now asks what the git pull command should do. The default option is recommended unless you specifically need to change its behavior. Click Next to continue with the installation.

17. Next you should choose which credential helper to use. Git uses credential helpers to fetch or save credentials. Leave the default option as it is the most stable one, and click Next.
18. select additional customization options and click next.
19. Once the installation is complete, tick the boxes to view the Release Notes or Launch Git Bash, then click Finish.

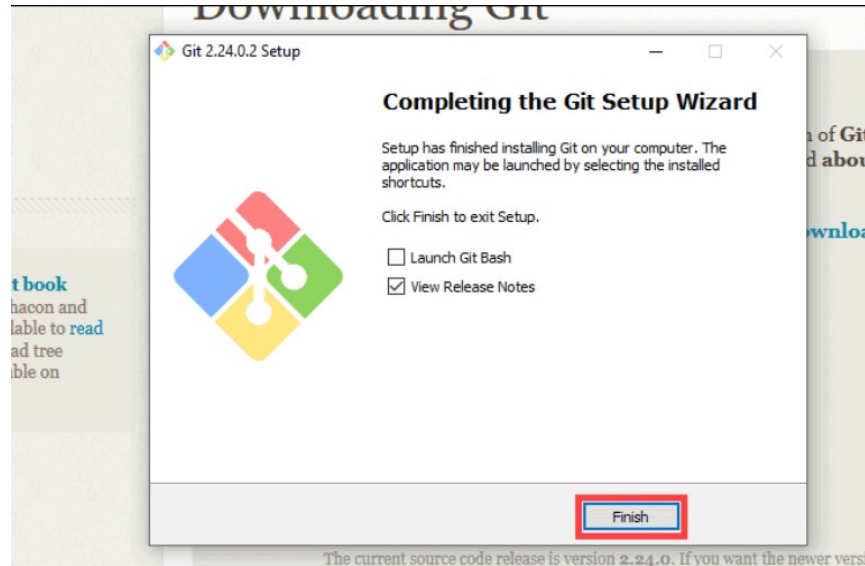


Figure 5: Completing the git set-up wizard

1.3 Install Git on Linux

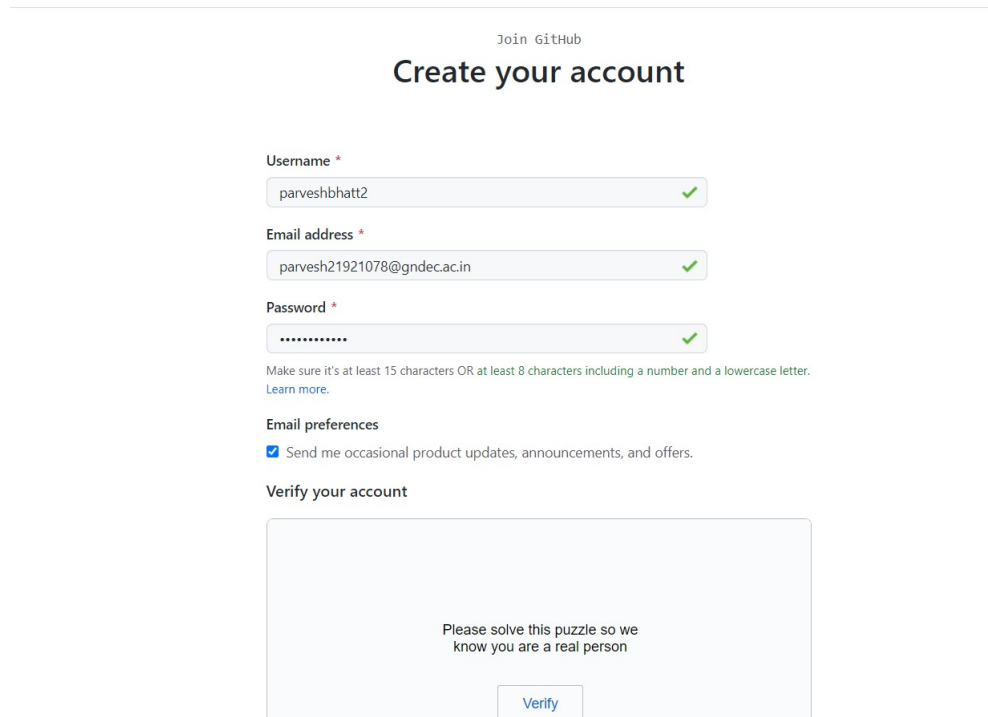
You can install Git on Linux through the package management tool that comes with your distribution.

1. Git packages are available using apt.
2. It's a good idea to make sure you're running the latest version. To do so, Navigate to your command prompt shell and run the following command to make sure everything is up-to-date: `sudo apt-get update`.
3. To install Git, run the following command: `sudo apt-get install git-all`.
4. Once the command output has completed, you can verify the installation by typing: `git version`.

2 PRACTICAL 2

CREATE AN ACCOUNT ON GITHUB

1. **Go to <https://github.com/join> in a web browser.** You can use any web browser on your computer, phone, or tablet to join.
2. **Enter your personal details.** In addition to creating a username and entering an email address, you'll also have to create a password. Your password must be at least 15 characters in length or at least 8 characters with at least one number and lower-case letter.



The screenshot shows the GitHub 'Create your account' page. At the top, it says 'Join GitHub' and 'Create your account'. Below this are three input fields: 'Username *' with the value 'parveshbhatt2', 'Email address *' with the value 'parvesh21921078@gndec.ac.in', and 'Password *' with masked characters. Each field has a green checkmark icon to its right. Below the password field, there is a note: 'Make sure it's at least 15 characters OR at least 8 characters including a number and a lowercase letter. [Learn more.](#)'. Underneath is the 'Email preferences' section with a checked checkbox for 'Send me occasional product updates, announcements, and offers.'. At the bottom, there is a 'Verify your account' section with a large box containing the text 'Please solve this puzzle so we know you are a real person' and a 'Verify' button.

Figure 6: Enter your personal details

3. **click the green Create an account button.** It's below the form.
4. **Complete the CAPTCHA puzzle.** The instructions vary by puzzle, so just follow the on screen instructions to confirm that you are a human.
5. **Click the Choose button for your desired plan.** Once you select a plan, GitHub will send an email confirmation message to the address you entered.
6. **Click the Verify email address button in the message from GitHub.** This confirms your email address and returns you to the sign up process.
7. **Review your plan selection and click Continue.** You can also choose whether you want to receive updates from GitHub via email by checking or unchecking the "Send me updates" box.
8. **your preferences and click Submit** GitHub displays a quick survey that can help you tailor your experience to match what you're looking for. Once you make your selection, you'll be taken to a screen that allows you to set up your first repository. Congratulations your account on GitHub is created Successfully. Now you are free to do your work.

3 PRACTICAL 3

CREATE REPOSITORY USING GIT/ GITHUB

GITHUB

1. In the upper-right corner of any page, use the drop-down menu, and select New repository.

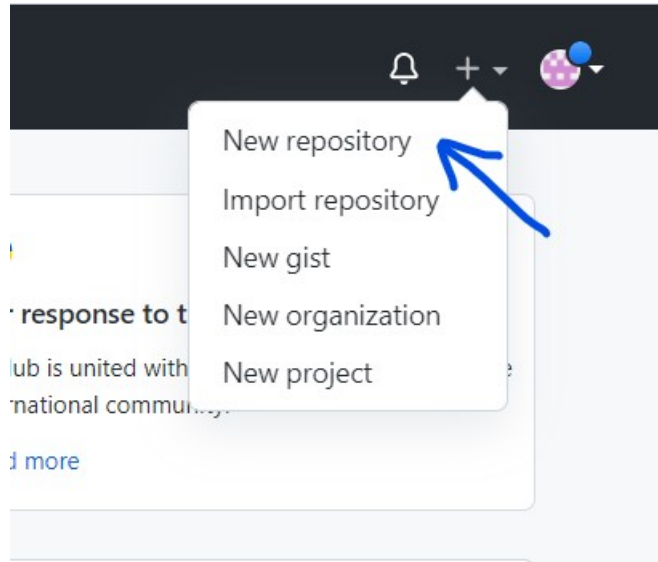



Figure 7: Select new repository

2. Type a short, memorable name for new repository. For example, "hello-devops".

Create a new repository


A repository contains all project files, including the revision history. Already have a project repository elsewhere? [Import a repository.](#)

Owner * Repository name *

 Parveshbhatt / hello-devops ✓

Great repository names are short and memorable. Need inspiration? How about [urban-octo-engine?](#)

Description (optional)

☒  **Public**
Anyone on the internet can see this repository. You choose who can commit.


☐  **Private**
You choose who can see and commit to this repository.

Figure 8: Name the repository

3. Optionally, add a description of our repository. For example, "First devops repository."

The screenshot shows the 'Name the repository' step in the GitHub repository creation process. It features two input fields: 'Owner' with a dropdown menu showing 'Parveshbhatt' and 'Repository name' with the text 'hello-devops' and a green checkmark. Below these is a text area for 'Description (optional)' containing 'First devops repository'. At the bottom, there are two radio button options: 'Public' (selected) and 'Private'. The 'Public' option is accompanied by a computer icon and the text 'Anyone on the internet can see this repository. You choose who can commit.' The 'Private' option is accompanied by a lock icon and the text 'You choose who can see and commit to this repository.'

Figure 9: Name the repository

4. Choose a repository visibility. For more information, see "About repositories".
5. Select Initialize this repository with a README.

The screenshot shows the 'Initialize this repository with' step. It includes a heading 'Initialize this repository with:' followed by the text 'Skip this step if you're importing an existing repository.' Below this are three checkboxes: 'Add a README file' (checked), 'Add .gitignore', and 'Choose a license'. Each checkbox has a brief description and a 'Learn more' link. At the bottom, there is a note: 'This will set `main` as the default branch. Change the default name in your settings.' A light blue information box states: 'You are creating a public repository in your personal account.' At the very bottom is a green button labeled 'Create repository'.

Figure 10: Select Add a README file

6. Click Create repository.

4 PRACTICAL 4

CREATE/DELETE/MERGE BRANCHES

4.1 CREATE BRANCH

A new branch can be created with the help of the *git branch* command.

```
$ git branch fix-code
```

where fix-code is the name of new branch.

4.2 MERGE BRANCH

Git allows you to merge the other branch with the currently active branch. You can merge two branches with the help of *git merge* command. Below command is used to merge the branches:

For merging a branch to master or any other branch, switch over to that branch using *git checkout* command.

```
$ git checkout master
```

Merge the fix-code branch to master/main branch.

```
$ git merge fix-code
```

The fix-code branch is now merged with master branch .

4.3 DELETE BRANCH

For deleting specified branch, it is a safe operation. In this command, Git prevents you from deleting the branch if it has unmerged changes. Below is the command to do this.

```
$ git branch -d fix-code
```

Attempting to delete an unmerged branch results in following error message.

```
error: The branch 'new' is not fully merged. If you are sure you want to delete it, run 'git branch -D new'.
```

To delete unmerged branch use same command with -D flag.

```
$ git branch -D branchname
```

For deleting a remote branch from Git desktop application. Below command is used to delete a remote branch:.

\$ git push origin -delete remote-branch

5 PRACTICAL 5

INSTALL DOCKER

Following are the steps to install Docker: -

1. Browse to official Docker website <https://docs.docker.com/desktop/windows/install/>
2. Double-click Docker Desktop Installer.exe to run the installer. If you haven't already downloaded the installer (Docker Desktop Installer.exe), you can get it from Docker Hub. It typically downloads to your Downloads folder, or you can run it from the recent downloads bar at the bottom of your web browser.



Figure 11: docker website

3. When prompted, ensure the Use WSL 2 instead of Hyper-V option on the Configuration page is selected or not depending on your choice of back-end.
4. Follow the instructions on the installation wizard to authorize the installer and proceed with the install.

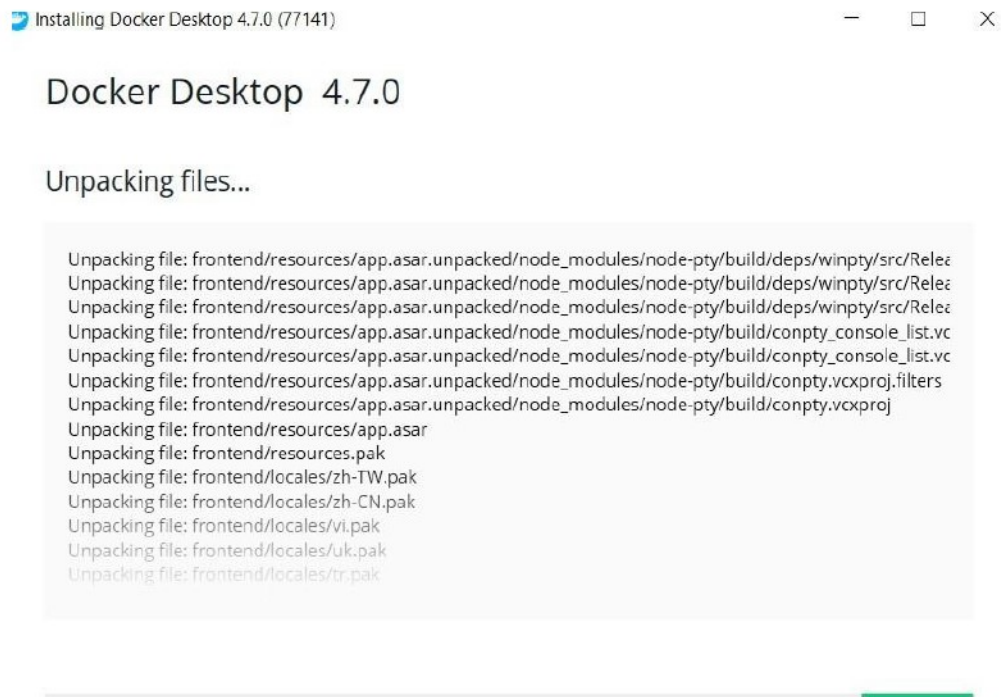


Figure 12: Installation wizard

- When the installation is successful, click Close to complete the installation process.




Figure 13: Installation Complete

- If your admin account is different to your user account, you must add the user to the docker-users group. Run Computer Management as an administrator and navigate to Local Users and Groups > Groups > dockerusers. Right-click to add the user to the group. Log out and log back in for the changes to take effect

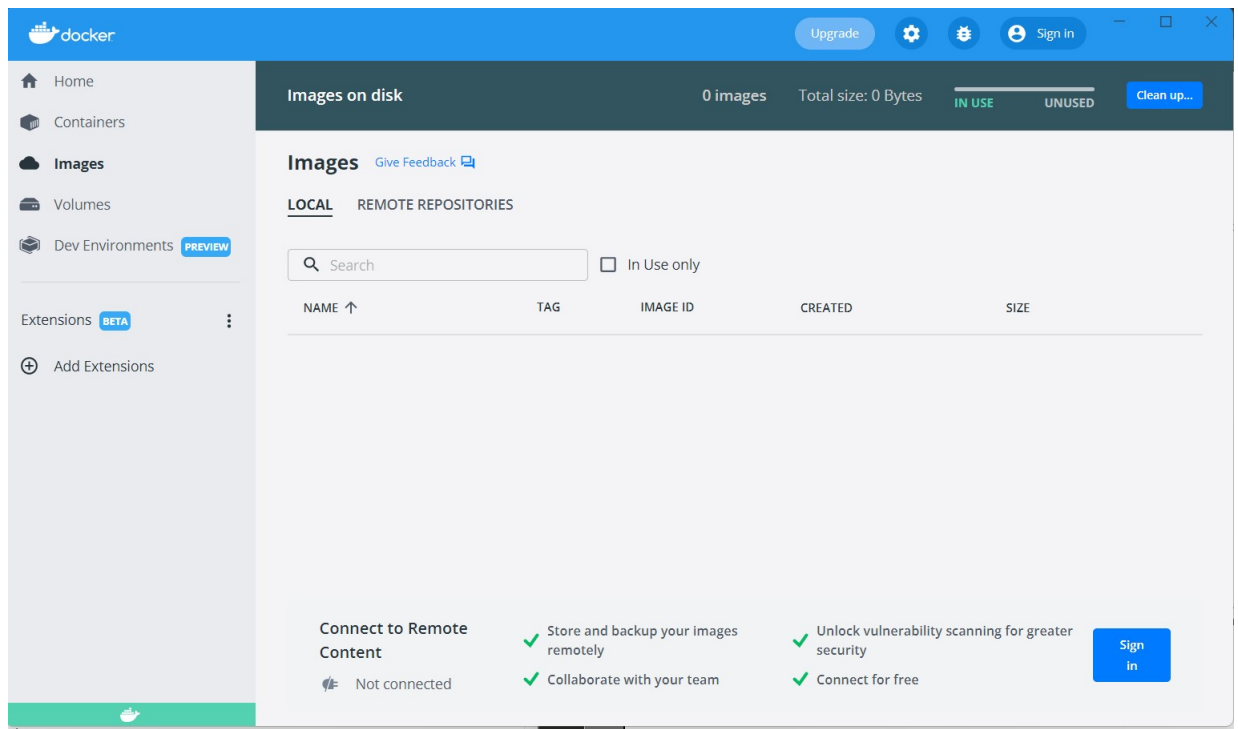


Figure 14: Docker Desktop

6 PRACTICAL 6

DEPLOY NGINX WEB SERVER IMAGE ON DOCKER

Steps for getting the Docker container for nginx up and running.

1. The first step is to pull the image from Docker Hub. When you log into Docker Hub, you will be able to search and see the image for nginx as shown below.

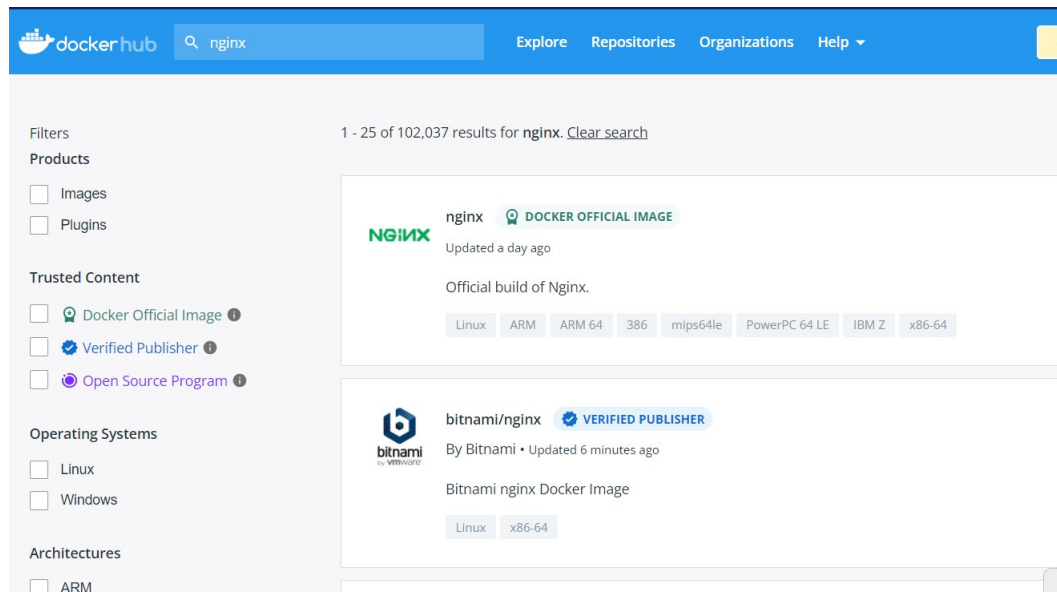


Figure 15: Nginx image on Docker hub

2. On the Docker Host, use the Docker pull command to download the latest nginx image from Docker Hub.

```
PS C:\Users\PARVESH> docker pull nginx
Using default tag: latest
latest: Pulling from library/nginx
214ca5fb9032: Pull complete
66eec13bb714: Pull complete
17cb812420e3: Pull complete
56fbf79cae7a: Pull complete
c4547ad15a20: Pull complete
d31373136b98: Pull complete
Digest: sha256:2d17cc4981bf1e22a87ef3b3dd20fbb72c3868738e3f307662eb40e2630d4320
Status: Downloaded newer image for nginx:latest
docker.io/library/nginx:latest
PS C:\Users\PARVESH> |
```

Figure 16: Pull nginx image

3. Now we have to run the NGINX image such that it will expose the Docker container port to the network port.

```
PS C:\Users\PARVESH> docker run -it --rm -d -p 8080:80 --name=webserver nginx
5793382f8ab1bff97b80a901d66f58605d6df66291caa601e63470f54a4d5bcb
PS C:\Users\PARVESH>
```

Figure 17: Run nginx image

Here, we have used the `-i` and `-t` options to run the container interactively, the `--rm` option to delete the container on exit, `-d` option to run the container in the background or detached mode, and the `--name` option to provide a name to the container. Finally, we have provided the name of the Docker image that we want to pull. We have also used the `-p` option to publish port 8080 of the container to port 80 of the host machine. Let's try to execute this command inside the terminal.

4. Let's verify by listing all the active containers.

```
PS C:\Users\PARVESH> docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS        PORTS                   NAMES
5793382f8ab1   nginx    "/docker-entrypoint..." 12 minutes ago Up 12 minutes  0.0.0.0:8080->80/tcp    webserver
PS C:\Users\PARVESH>
```

Figure 18: Listing running containers

5. To see that the container is actively running, navigate to the link `localhost:8080` inside a browser, we will find the welcome page of the Nginx web server.

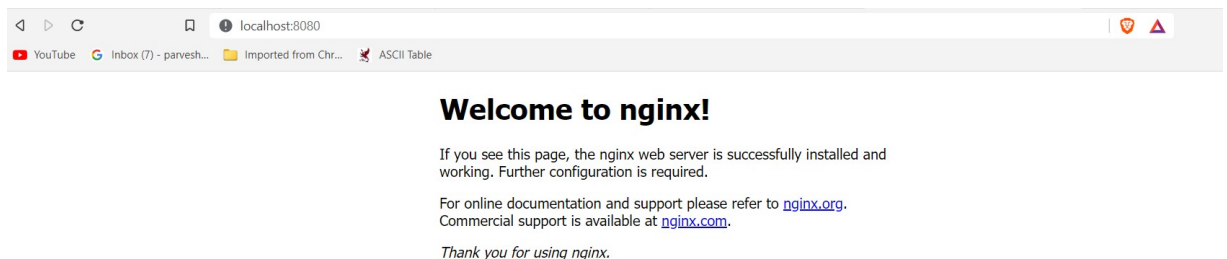


Figure 19: Nginx welcome page

6. To stop the container use the 'docker stop' command


```
PS C:\Users\PARVESH> docker stop webserver
webserver
PS C:\Users\PARVESH> docker ps
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS        NAMES
PS C:\Users\PARVESH> docker ps -a
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS        NAMES
1310d5af0e11   feb5d9fea6a5  "/hello"                12 days ago   Exited (0)    12 days ago   charming_morse
PS C:\Users\PARVESH> |
```

Figure 20: Stop Nginx

7 PRACTICAL 7

DEPLOY APACHE WEB SERVER IMAGE ON DOCKER

Steps for getting the Docker container for Apache Web Server Image on Docker.

1. Log into Docker Hub, you will be able to search and see the image for apache as httpd.

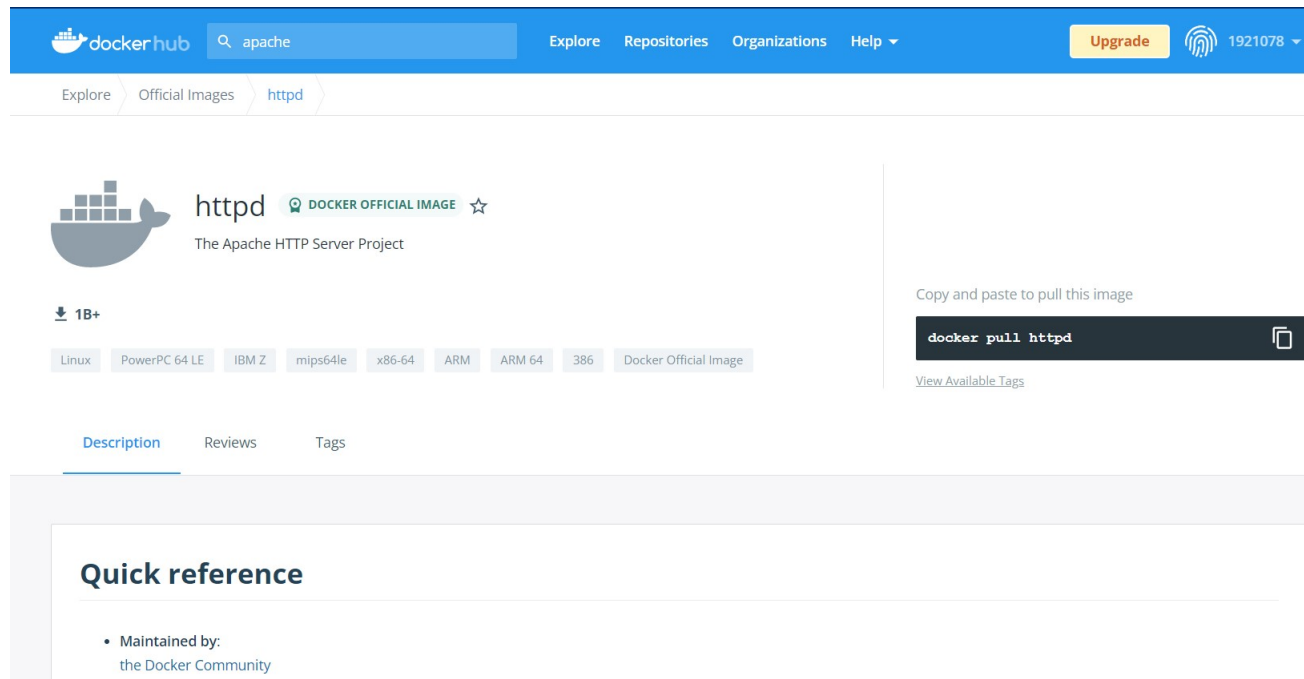


Figure 21: httpd Image On Docker Hub

2. Pull the Docker image, which contains Apache called httpd, by running the docker pull command below. This command will download or pull the Apache image from the Docker registry, as shown below.

```
PS C:\Users\PARVESH> docker pull httpd
Using default tag: latest
latest: Pulling from library/httpd
214ca5fb9032: Already exists
7cf31a2ecec6: Pull complete
bf666e57b9f2: Pull complete
c15a4e94ae6b: Pull complete
dc25474c7f97: Pull complete
Digest: sha256:2d1f8839d6127e400ac5f65481d8a0f17ac46a3b91de40b01e649c9a0324dea0
Status: Downloaded newer image for httpd:latest
docker.io/library/httpd:latest
PS C:\Users\PARVESH>
```

Figure 22: Pull httpd

3. Next, confirm the downloaded image by running the docker images command below to list all images available on your computer.

```
PS C:\Users\PARVESH> docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
nginx         latest    de2543b9436b   7 days ago    142MB
httpd         latest    c58ef9bfbb57   2 weeks ago    144MB
PS C:\Users\PARVESH>
```


Figure 23: Check Images

4. Invoke the docker run command to create a new container based on your downloaded Apache Docker image. The docker run command then returns the unique container ID of the container you've just created. Save this container id in the highlighted box below to the future if you wish to delete or remove the container.

```
PS C:\Users\PARVESH> docker run -dit --name my-apache-app -p 8080:80 httpd
8dec14cc3c33bab786c2d1f46521a1cb07f4f31103ddb5db88162b97e88e7a
PS C:\Users\PARVESH>
```

Figure 24: Run Server

5. Once the Apache container is running, verify it by access the Apache web interface by navigating to Public-Ip-address:80 using any web browser. If you can see the same message, as you can see below, then you've successfully started your Apache Docker container.



It works!

Figure 25: Default output of Apache web server

8 PRACTICAL 8

CREATE CUSTOM PAGE USING WEB SERVER

1. First we create a running container. The docker create command will create a new container. Here we have requested a new container named baseimg with port 80 exposed to localhost. We are using nginx as a base image for the container. If you don't have the nginx image in your local docker image repository, it will download automatically.

```
PS C:\Users\PARVESH> docker create --name base_img -p 8080:80 nginx
9f14980e1684f701e820ac4b5e225f423907f394acd48e8c136cdf5ee0639501
PS C:\Users\PARVESH> |
```

Figure 26: Creating a base container

2. If you look at the list of images on your system, you will now see the nginx:alpine image.

```
PS C:\Users\PARVESH> docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
nginx         latest    de2543b9436b   7 days ago    142MB
httpd         latest    c58ef9bfbb57   2 weeks ago   144MB
PS C:\Users\PARVESH>
```

Figure 27: Inspect Images

3. Note here that the container is not running, so you won't see it in the container list unless you use the -a flag (-a is for all).

```
PS C:\Users\PARVESH> docker ps -a
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS              PORTS          NAMES
9f14980e1684   nginx     "/docker-entrypoint..." 3 minutes ago   Created                                base_img
8dec14cc3c33   httpd     "httpd-foreground"       44 minutes ago   Exited (0) 27 minutes ago              my-apache-app
1310d5af0e11   feeb5d9fea6a5 "/hello"                13 days ago     Exited (0) 13 days ago              charming_morse
PS C:\Users\PARVESH>
```


Figure 28: Inspect Container

4. Start the container by running the following command.

```
PS C:\Users\PARVESH> docker start base_img
base_img
PS C:\Users\PARVESH> |
```

Figure 29: Start the container

Now visit <http://localhost> with your browser. You will see the default “Welcome to nginx!” page. We are now running an nginx container.



It works!

Figure 30: Container Running

5. Let's create a new index.html file and copy it onto the running container. Using an editor on your machine, create an index.html file in the same directory that you have been running Docker commands from. Then paste the following HTML into it:

```
<html>
  <head>
    <title>Hello Page</title>
  </head>
  <body>
    <h1>Hello world</h1>
  </body>
</html>
```

Figure 31: Html File

6. Then save the file and return to the command line. We will use the docker cp command to copy this file onto the running container.

```
PS C:\Users\PARVESH> docker cp "D:\index.html" base_img:/usr/share/nginx/html/index.html
PS C:\Users\PARVESH> |
```

Figure 32: Copy file on running container

7. Now reload your browser or revisit <http://localhost>. You will see the message “Hello World!” in place of the default nginx welcome page.




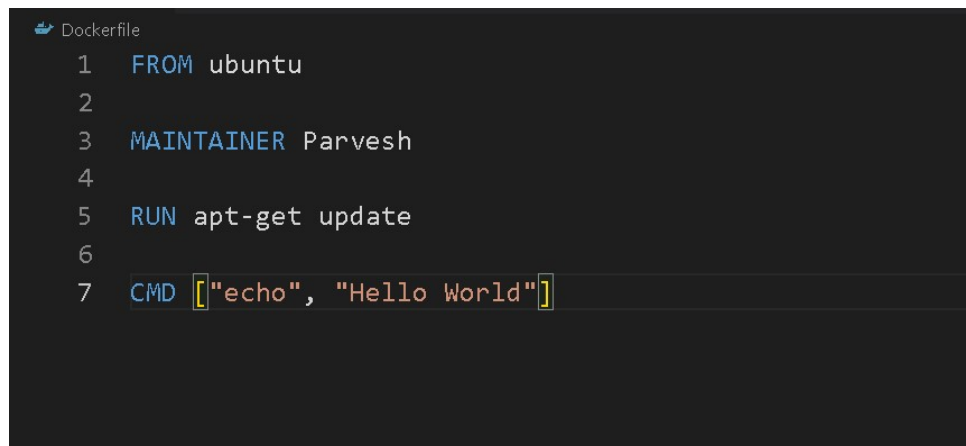
Figure 33: Modified page

9 PRACTICAL 9

CREATE CUSTOM IMAGE

Steps to Create a custom image

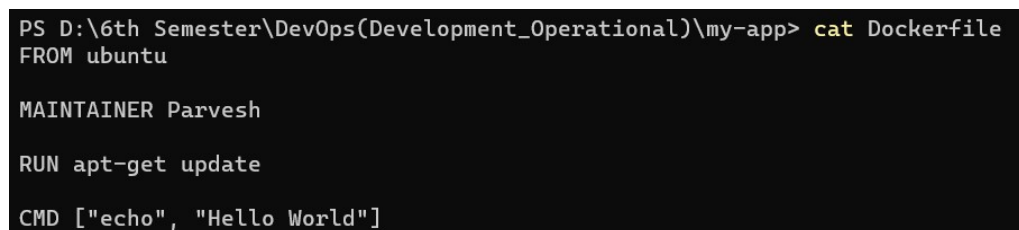
1. First we have to create a directory named my-app and move into that directory and create a new empty file (Dockerfile).
2. write the following commands in Dockerfile.



```
Dockerfile
1 FROM ubuntu
2
3 MAINTAINER Parvesh
4
5 RUN apt-get update
6
7 CMD ["echo", "Hello World"]
```

Figure 34: Dockerfile

3. Save and exit the file.
4. You can check the content of the file by using the cat command.



```
PS D:\6th Semester\DevOps(Development_Operational)\my-app> cat Dockerfile
FROM ubuntu

MAINTAINER Parvesh

RUN apt-get update

CMD ["echo", "Hello World"]
```

Figure 35: Checking content of Dockerfile

5. The basic syntax used to build an image using a Dockerfile is: ***docker build [OPTIONS] PATH — URL —***
6. Build a docker image, If you are already in the directory where the Dockerfile is located, put a . instead of the location. By adding the -t flag, you can tag the new image with a name which will help you when dealing with multiple images.

```

PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker build -t myapp:1.0 .
[+] Building 132.3s (6/6) FINISHED
=> [internal] load build definition from Dockerfile                                0.6s
=> => transferring dockerfile: 123B                                              0.5s
=> [internal] load .dockerignore                                                  0.4s
=> => transferring context: 2B                                                  0.3s
=> [internal] load metadata for docker.io/library/ubuntu:latest                 10.9s
=> [1/2] FROM docker.io/library/ubuntu@sha256:b6b83d3c331794420340093eb706a6f152d9c1fa51b262d9bf34594887c2c7ac 16.0s
=> => resolve docker.io/library/ubuntu@sha256:b6b83d3c331794420340093eb706a6f152d9c1fa51b262d9bf34594887c2c7ac 0.0s
=> => sha256:b6b83d3c331794420340093eb706a6f152d9c1fa51b262d9bf34594887c2c7ac 1.42kB / 1.42kB 0.0s
=> => sha256:bace9fb0d5923a675c894d5c815da75ffe35e24970166a48a4460a48ae6e0d19 529B / 529B 0.0s
=> => sha256:27941809078cc9b2802deb2b0bb6feed6c236cde01e487f200e24653533701ee 1.46kB / 1.46kB 0.0s
=> => sha256:405f018f9d1d0f351c196b841a7c7f226fb8ea448acd6339a9ed8741600275a2 30.42MB / 30.42MB 7.3s
=> => extracting sha256:405f018f9d1d0f351c196b841a7c7f226fb8ea448acd6339a9ed8741600275a2 6.3s
=> [2/2] RUN apt-get update                                                    103.4s
=> => exporting to image                                                         1.0s
=> => exporting layers                                                         0.9s
=> => writing image sha256:92ffb180c683dd1fbd65c675187d80aa85b04de61f979ccbb544680e76450988 0.0s
=> => naming to docker.io/library/myapp:1.0                                    0.0s

Use 'docker scan' to run Snyk tests against images to find vulnerabilities and learn how to fix them

```

Figure 36: Building Image

- Once the image is successfully built, you can verify whether it is on the list of local images with the command.

```

PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker images
REPOSITORY    TAG       IMAGE ID       CREATED        SIZE
myapp         1.0      92ffb180c683  About a minute ago  112MB
nginx         latest   de2543b9436b  3 weeks ago    142MB
httpd         latest   c58ef9bfb57   4 weeks ago    144MB

```

Figure 37: Verify image

- Launch a new Docker container based on the image you created in the previous steps.

```

PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker run myapp
Unable to find image 'myapp:latest' locally
docker: Error response from daemon: pull access denied for myapp, repository does not exist or may require 'docker login': denied: requested access to the resource is denied.
See 'docker run --help'.
PS D:\6th Semester\DevOps(Development_Operational)\my-app> d
d: The term 'd' is not recognized as a name of a cmdlet, function, script file, or executable program.
Check the spelling of the name, or if a path was included, verify that the path is correct and try again.
PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker run myapp:1.0
Hello World

```

Figure 38: Launching Docker container

10 PRACTICAL 10

PUSH CUSTOM IMAGE TO DOCKER HUB

Steps to push a custom image to Docker hub are: -

1. Docker images are pushed to Docker Hub through the docker push command. A single Docker Hub repository can hold many Docker images (stored as tags).
2. To create a repository, sign into Docker Hub, click on Repositories then Create Repository.

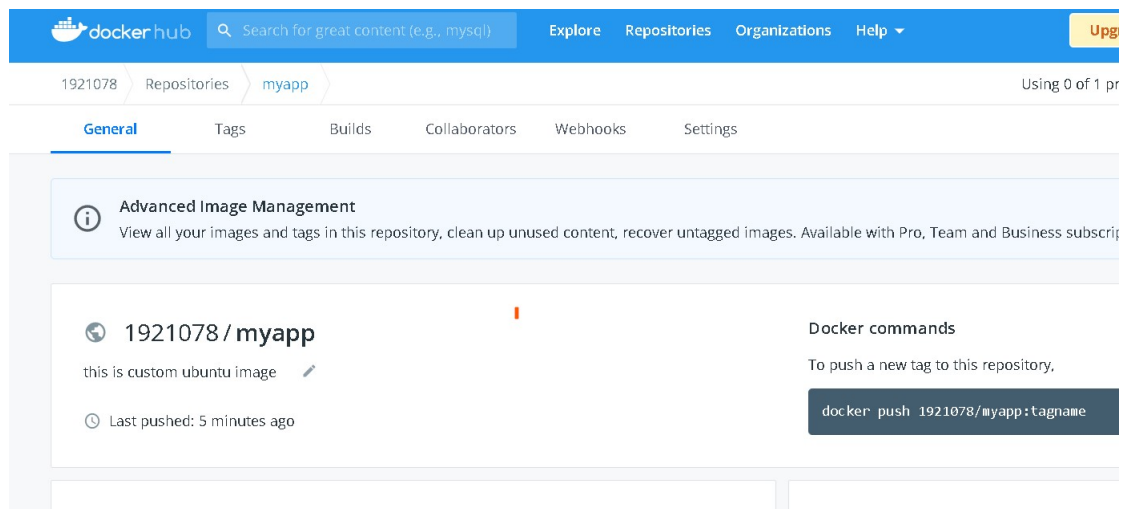


Figure 39: Create Repository

3. To push an image to Docker Hub, you must first name your local image using your Docker Hub username and the repository name that you created through Docker Hub on the web.
4. Name your local images using re-tagging an existing local image *dockertag < existing-image >< hub - user > / < repo - name > [:< tag >]*

```
PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker tag myapp:1.0 1921078/myapp:1.0
PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker images
REPOSITORY          TAG                 IMAGE ID            CREATED             SIZE
1921078/myapp        1.0                92ffb180c683       57 minutes ago     112MB
myapp                1.0                92ffb180c683       57 minutes ago     112MB
nginx                latest             de2543b9436b       3 weeks ago        142MB
httpd                latest             c58ef9bfbb57       4 weeks ago        144MB
```

Figure 40: Re-tagging Repository

5. Push this repository to the registry designated by its name or tag.

```
PS D:\6th Semester\DevOps(Development_Operational)\my-app> docker push 1921078/myapp:1.0
The push refers to repository [docker.io/1921078/myapp]
541285584b20: Pushed
a790f937a6ae: Pushed
1.0: digest: sha256:ad1d8d05c0093c4f0bf6a9d0d40d7325dc8d5d019165dbc71f719c2f6d05f93d size: 741
```

Figure 41: Push repository

- The image is then uploaded and available for use by your teammates and/or the community.

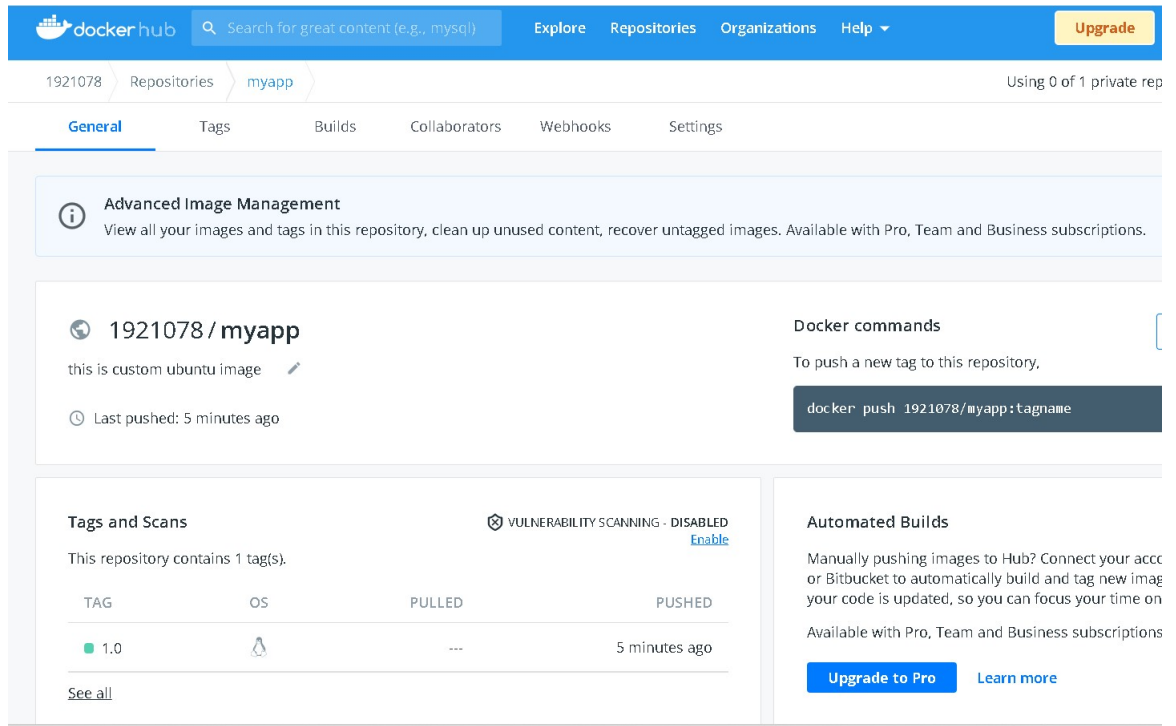


Figure 42: repository(Image) on Docker hub

11 PRACTICAL 11

USE PERSISTENT STORAGE WITH DOCKER

Docker has two options for containers to store files on the host machine, so that the files are persisted even after the container stops: volumes, and bind mounts.

Docker also supports containers storing files in-memory on the the host machine. Such files are not persisted. If you're running Docker on Linux, tmpfs mount is used to store files in the host's system memory. If you're running Docker on Windows, named pipe is used to store files in the host's system memory.

1. Volumes are stored in a part of the host filesystem which is managed by Docker (/var/lib/docker/volumes/ on Linux). Non-Docker processes should not modify this part of the filesystem. Volumes are the best way to persist data in Docker.
2. Bind mounts may be stored anywhere on the host system. They may even be important system files or directories. Non-Docker processes on the Docker host or a Docker container can modify them at any time.

Steps to use Docker Volumes are: -

1. There are different ways to mount a Docker volume while launching a container. Users can decide between the -v and the --mount flags, which are added to the docker run command.
2. create a Docker volume.

```
PS C:\Users\PARVESH> docker volume create mydata
mydata
```

Figure 43: Create Volume

3. Verify you have successfully created a Docker volume, prompt Docker to list all available volumes.

```
PS C:\Users\PARVESH> docker volume ls
DRIVER      VOLUME NAME
local       mydata
```

Figure 44: List volumes

4. To see more information about a Docker volume, use the inspect command.

```
PS C:\Users\PARVESH> docker inspect mydata
[
  {
    "CreatedAt": "2022-06-10T06:35:13Z",
    "Driver": "local",
    "Labels": {},
    "Mountpoint": "/var/lib/docker/volumes/mydata/_data",
    "Name": "mydata",
    "Options": {},
    "Scope": "local"
  }
]
```

Figure 45: Inspect volume

5. To mount a data volume to a container add the `--mount` flag to the docker run command. It adds the volume to the specified container, where it stores the data produced inside the virtual environment.

To run a container and mount a data volume to it, follow the basic syntax:

$dockerrun - mountsource = [volume_name], destination = [path_in_container][docker_image]$

6. To launch an Ubuntu container and mount the mydata volume to it, run:

```
PS C:\Users\PARVESH> docker run -it --name=example1 --mount source=mydata,destination=/data ubuntu
Unable to find image 'ubuntu:latest' locally
latest: Pulling from library/ubuntu
405f018f9d1d: Already exists
Digest: sha256:b6b83d3c331794420340093eb706a6f152d9c1fa51b262d9bf34594887c2c7ac
Status: Downloaded newer image for ubuntu:latest
```

Figure 46: Mount mydata volume

The steps to Copying Files Between Containers From a Shared Volume are: -

- (a) switched to the container command prompt, move to the data volume directory.
- (b) Create an empty sample file using the touch command.
- (c) Now, exit the container.

```
root@bfc8f5bb198f:/# ls
bin  data  etc  lib  lib64  media  opt  root  sbin  sys  usr
boot dev  home lib32 libx32 mnt  proc  run  srv  tmp  var
root@bfc8f5bb198f:/# cd data
root@bfc8f5bb198f:/data# ls
root@bfc8f5bb198f:/data# touch sample1.txt
root@bfc8f5bb198f:/data# ls
sample1.txt
root@bfc8f5bb198f:/data# exit
exit
```

Figure 47: File in first container

- (d) Then, launch a new container example2 with the same data volume.

```
PS C:\Users\PARVESH> docker run -it --name=example2 --mount source=mydata,destination=/data ubuntu
root@bfc8f5bb198f:/#
```

Figure 48: Launch example2 Container

- (e) List the content of the container. You should find the data directory, as in example1.
- (f) Move to the data directory and list the content of it.
The output should list the sample1.txt file you created in the previous container (example1).

```

root@d51afdb14705:/# ls
bin  dev  lib  libx32  opt  run  sys  var
boot  etc  lib32  media  proc  sbin  tmp
data  home  lib64  mnt     root  srv   usr
root@d51afdb14705:/# cd data
root@d51afdb14705:/data# ls
sample1.txt
root@d51afdb14705:/data# exit
exit
PS C:\Users\PARVESH>

```

Figure 49: File in second container

You can also mount an existing directory from the host machine to a container. This type of volume is called Host Volumes.

You can mount host volumes by using the `-v` flag and specifying the name of the host directory.

The basic syntax for mounting a host directory is:

`dockerrun -v "$(pwd)" : [volume_name][docker_image]`

The `"$pwd"` attribute instructs Docker to mount the directory the user is currently in. The following example outlines how this is done.

1. First, create a sample directory on the host under the name `tmp` and move into it:
2. Once inside the directory, create a test file to see whether it will be available from the container.

```

PS D:\6th Semester\DevOps(Development_Operational)> cd tmp
PS D:\6th Semester\DevOps(Development_Operational)\tmp> ls

Directory: D:\6th Semester\DevOps(Development_Operational)\tmp

Mode                LastWriteTime         Length Name
----                -
-a---             10-06-2022   12:54             0 file1.txt

```

Figure 50: Sample Directory

3. Then, use the `docker run` command to launch an Ubuntu container with the host directory attached to it.
4. List the content of the container and verify there is a `data1` directory.
5. Open the mounted directory and list the content. The output should display the file you created on the host.

```
PS D:\6th Semester\DevOps(Development_Operational)\tmp> docker run -it -v ${
PWD}:/data1 ubuntu
root@35b165319d40:/# ls
bin  data1  etc  lib  lib64  media  opt  root  sbin  sys  usr
boot  dev  home  lib32  libx32  mnt  proc  run  srv  tmp  var
root@35b165319d40:/# cd data1
root@35b165319d40:/data1# ls
file1.txt
```

Figure 51: File in the container